

```

1 # Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py
2
3 %matplotlib inline
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6
7 import warnings
8 warnings.filterwarnings("ignore")
9
10 from __future__ import print_function
11 import keras
12 from keras.datasets import mnist
13 from keras.models import Sequential
14 from keras.layers import Dense, Dropout, Flatten
15 from keras.layers import Conv2D, MaxPooling2D
16 from keras import backend as K
17
18 batch_size = 128
19 num_classes = 10
20 epochs = 10
21
22 # input image dimensions
23 img_rows, img_cols = 28, 28
24
25 # the data, split between train and test sets
26 (x_train, y_train), (x_test, y_test) = mnist.load_data()
27
28
29 if K.image_data_format() == 'channels_first':
30     x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
31     x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
32     input_shape = (1, img_rows, img_cols)
33 else:
34     x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
35     x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
36     input_shape = (img_rows, img_cols, 1)
37
38 x_train = x_train.astype('float32')
39 x_test = x_test.astype('float32')
40 x_train /= 255
41 x_test /= 255
42 print('x_train shape:', x_train.shape)
43 print(x_train.shape[0], 'train samples')
44 print(x_test.shape[0], 'test samples')
45
46 # convert class vectors to binary class matrices
47 y_train = keras.utils.to_categorical(y_train, num_classes)
48 y_test = keras.utils.to_categorical(y_test, num_classes)
49

```

```

↳ x_train shape: (60000, 28, 28, 1)
   60000 train samples
   10000 test samples

```

```

1 import numpy as np
2 import time
3 # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
4 # https://stackoverflow.com/a/14434334
5 # this function is used to update the plots for each epoch and error
6 def plt_dynamic(x, vy, ty, fig, ax, colors=['b']):
7     ax.plot(x, vy, 'b', label="Validation Loss")
8     ax.plot(x, ty, 'r', label="Train Loss")
9     plt.legend()
10    plt.grid()
11    fig.canvas.draw()

```

```

1 def loss_plot(model_name):
2     score = model_name.evaluate(x_test, y_test, verbose=0)
3     print('Test score:', score[0])
4     print('Test accuracy:', score[1])
5
6     fig, ax = plt.subplots(1,1, figsize=(10,6))
7     ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
8     ax.set_title('Variation of Loss with epochs')
9
10    # list of epoch numbers
11    x = list(range(1,epochs+1))
12
13    # print(history.history.keys())
14    # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
15    # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
16
17    # we will get val_loss and val_acc only when you pass the paramter validation_data
18    # val_loss : validation loss
19    # val_acc : validation accuracy
20
21    # loss : training loss
22    # acc : train accuracy
23    # for each key in history.history we will have a list of length equal to number of epochs
24
25
26    vy = history.history['val_loss']
27    ty = history.history['loss']
28    plt_dynamic(x, vy, ty, fig, ax)

```

```

1 def weight_plot(model_name):
2     w_after = model_name.get_weights()
3
4     h1_w = w_after[0].flatten().reshape(-1,1)

```

```

5     h2_w = w_after[2].flatten().reshape(-1,1)
6     out_w = w_after[4].flatten().reshape(-1,1)
7
8     fig = plt.figure(figsize=(15,7))
9     fig.suptitle("Weight matrices after model trained")
10    plt.subplot(1, 3, 1)
11    ax = sns.violinplot(y=h1_w,color='b')
12    plt.xlabel('Hidden Layer 1')
13
14    plt.subplot(1, 3, 2)
15    ax = sns.violinplot(y=h2_w, color='r')
16    plt.xlabel('Hidden Layer 2 ')
17
18    plt.subplot(1, 3, 3)
19    ax = sns.violinplot(y=out_w,color='y')
20    plt.xlabel('Output Layer ')
21    plt.show()

```

▼ Architecture 1:

no. of Layers = 3; Kernel_size = 3x3

```

1  model1 = Sequential()
2  #Layer 1
3  model1.add(Conv2D(48, kernel_size=(3, 3),
4                    activation='relu',
5                    input_shape=input_shape))
6  #Layer 2
7  model1.add(Conv2D(64, (3, 3), activation='relu'))
8  model1.add(MaxPooling2D(pool_size=(2, 2)))
9  model1.add(Dropout(0.5))
10
11 #Layer 3
12 model1.add(Conv2D(50, (3, 3), activation='relu'))
13 model1.add(MaxPooling2D(pool_size=(2, 2)))
14 model1.add(Dropout(0.5))
15
16 model1.add(Flatten())
17 model1.add(Dense(128, activation='relu'))
18 model1.add(Dropout(0.5))
19 model1.add(Dense(num_classes, activation='softmax'))
20 print(model1.summary())
21
22 model1.compile(loss=keras.losses.categorical_crossentropy,
23               optimizer=keras.optimizers.Adadelta(),
24               metrics=['accuracy'])
25
26 history = model1.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
27

```



Model: "sequential_5"



Layer (type)	Output Shape	Param #
conv2d_18 (Conv2D)	(None, 26, 26, 48)	480
conv2d_19 (Conv2D)	(None, 24, 24, 64)	27712
max_pooling2d_14 (MaxPooling)	(None, 12, 12, 64)	0
dropout_16 (Dropout)	(None, 12, 12, 64)	0
conv2d_20 (Conv2D)	(None, 10, 10, 50)	28850
max_pooling2d_15 (MaxPooling)	(None, 5, 5, 50)	0
dropout_17 (Dropout)	(None, 5, 5, 50)	0
flatten_4 (Flatten)	(None, 1250)	0
dense_7 (Dense)	(None, 128)	160128
dropout_18 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 10)	1290
Total params: 218,460		
Trainable params: 218,460		
Non-trainable params: 0		

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 5s 88us/step - loss: 0.3946 - acc: 0.8724 - val_loss: 0.0837 - val_acc: 0.9743
Epoch 2/10
60000/60000 [=====] - 4s 71us/step - loss: 0.1470 - acc: 0.9553 - val_loss: 0.0469 - val_acc: 0.9843
Epoch 3/10
60000/60000 [=====] - 4s 72us/step - loss: 0.1110 - acc: 0.9662 - val_loss: 0.0352 - val_acc: 0.9877
Epoch 4/10
60000/60000 [=====] - 4s 72us/step - loss: 0.0968 - acc: 0.9710 - val_loss: 0.0309 - val_acc: 0.9887
Epoch 5/10
60000/60000 [=====] - 4s 72us/step - loss: 0.0848 - acc: 0.9741 - val_loss: 0.0302 - val_acc: 0.9901
Epoch 6/10
60000/60000 [=====] - 4s 72us/step - loss: 0.0780 - acc: 0.9770 - val_loss: 0.0260 - val_acc: 0.9916
Epoch 7/10
60000/60000 [=====] - 4s 72us/step - loss: 0.0710 - acc: 0.9791 - val_loss: 0.0257 - val_acc: 0.9910
Epoch 8/10
60000/60000 [=====] - 4s 72us/step - loss: 0.0669 - acc: 0.9801 - val_loss: 0.0243 - val_acc: 0.9916
Epoch 9/10
60000/60000 [=====] - 4s 72us/step - loss: 0.0658 - acc: 0.9804 - val_loss: 0.0214 - val_acc: 0.9940
Epoch 10/10
60000/60000 [=====] - 4s 72us/step - loss: 0.0623 - acc: 0.9818 - val_loss: 0.0207 - val_acc: 0.9935

```
1 loss_plot(model1)
2 weight_plot(model1)
```



Test score: 0.02070148353522236
Test accuracy: 0.9935



Architecture 2:

no. of Layers = 5; Kernel_size = 5x5

```
1 model2 = Sequential()
2 #Layer 1
3 model2.add(Conv2D(50, kernel_size=(5, 5), padding='same',
4                   activation='relu',
5                   input_shape=input_shape))
6
7 #Layer 2
8 model2.add(Conv2D(50, (5, 5), padding='same', activation='relu'))
9 model2.add(MaxPooling2D(pool_size=(2, 2)))
10 model2.add(Dropout(0.5))
11
12 #Layer 3
13 model2.add(Conv2D(60, (5, 5), padding='same', activation='relu'))
14 model2.add(Dropout(0.5))
15
16 #Layer 4
17 model2.add(Conv2D(40, (5, 5), padding='same', activation='relu'))
18 model2.add(MaxPooling2D(pool_size=(3, 3), padding='same'))
19 model2.add(Dropout(0.5))
20
21 #Layer 5
22 model2.add(Conv2D(30, (5, 5), padding='same', activation='relu'))
23 model2.add(Dropout(0.5))
24
25 model2.add(Flatten())
26 model2.add(Dense(128, activation='relu'))
27 model2.add(Dropout(0.5))
28 model2.add(Dense(num_classes, activation='softmax'))
29
30 print(model2.summary())
31
32 model2.compile(loss=keras.losses.categorical_crossentropy,
```

```

33         optimizer=keras.optimizers.Adadelta(),
34         metrics=['accuracy'])
35
36 history = model2.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
37

```

```

↳ Layer (type)                Output Shape                Param #
=====
conv2d_70 (Conv2D)            (None, 28, 28, 50)         1300
conv2d_71 (Conv2D)            (None, 28, 28, 50)         62550
max_pooling2d_55 (MaxPooling (None, 14, 14, 50)         0
dropout_70 (Dropout)          (None, 14, 14, 50)         0
conv2d_72 (Conv2D)            (None, 14, 14, 60)         75060
dropout_71 (Dropout)          (None, 14, 14, 60)         0
conv2d_73 (Conv2D)            (None, 14, 14, 40)         60040
max_pooling2d_56 (MaxPooling (None, 5, 5, 40)         0
dropout_72 (Dropout)          (None, 5, 5, 40)         0
conv2d_74 (Conv2D)            (None, 5, 5, 30)          30030
dropout_73 (Dropout)          (None, 5, 5, 30)         0
flatten_12 (Flatten)          (None, 750)                0
dense_23 (Dense)              (None, 128)                96128
dropout_74 (Dropout)          (None, 128)                0
dense_24 (Dense)              (None, 10)                 1290
=====
Total params: 326,398
Trainable params: 326,398
Non-trainable params: 0

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 11s 189us/step - loss: 0.6108 - acc: 0.7979 - val_loss: 0.0709 - val_acc: 0.97
Epoch 2/10
60000/60000 [=====] - 7s 112us/step - loss: 0.1488 - acc: 0.9578 - val_loss: 0.0718 - val_acc: 0.977
Epoch 3/10
60000/60000 [=====] - 7s 112us/step - loss: 0.1041 - acc: 0.9701 - val_loss: 0.0332 - val_acc: 0.989
Epoch 4/10
60000/60000 [=====] - 7s 112us/step - loss: 0.0875 - acc: 0.9750 - val_loss: 0.0368 - val_acc: 0.988
Epoch 5/10
60000/60000 [=====] - 7s 114us/step - loss: 0.0779 - acc: 0.9792 - val_loss: 0.0254 - val_acc: 0.991
Epoch 6/10
60000/60000 [=====] - 7s 113us/step - loss: 0.0669 - acc: 0.9810 - val_loss: 0.0226 - val_acc: 0.993
Epoch 7/10
60000/60000 [=====] - 7s 113us/step - loss: 0.0606 - acc: 0.9826 - val_loss: 0.0211 - val_acc: 0.992
Epoch 8/10
60000/60000 [=====] - 7s 114us/step - loss: 0.0603 - acc: 0.9838 - val_loss: 0.0199 - val_acc: 0.994
Epoch 9/10
60000/60000 [=====] - 7s 116us/step - loss: 0.0566 - acc: 0.9851 - val_loss: 0.0303 - val_acc: 0.990
Epoch 10/10
60000/60000 [=====] - 7s 112us/step - loss: 0.0511 - acc: 0.9860 - val_loss: 0.0194 - val_acc: 0.994

```

```

1 loss_plot(model2)


```

```

↳

```

Test score: 0.019397488478862943
Test accuracy: 0.994



▼ Architecture 3:

no. of Layers = 7; Kernel_size = 6x6

along with Batch Normalization and Dropout layers

```
1 from keras.layers.normalization import BatchNormalization
2
3 model3 = Sequential()
4 #Layer 1
5 model3.add(Conv2D(70, kernel_size=(6, 6), padding='same',
6 activation='relu',
7 input_shape=input_shape))
8 model3.add(BatchNormalization())
9
10 #Layer 2
11 model3.add(Conv2D(80, (6, 6), padding='same', activation='relu'))
12 model3.add(BatchNormalization())
13 model3.add(MaxPooling2D(pool_size=(2, 2), padding='same'))
14 model3.add(Dropout(0.5))
15
16 #Layer 3
17 model3.add(Conv2D(60, (6, 6), padding='same', activation='relu'))
18 model3.add(BatchNormalization())
19 model3.add(Dropout(0.5))
20
21 #Layer 4
22 model3.add(Conv2D(50, (6, 6), padding='same', activation='relu'))
23 model3.add(BatchNormalization())
24 model3.add(MaxPooling2D(pool_size=(2, 2), padding='same'))
25 model3.add(Dropout(0.5))
26
27 #Layer 5
28 model3.add(Conv2D(48, (6, 6), padding='same', activation='relu'))
29 model3.add(BatchNormalization())
30 model3.add(Dropout(0.5))
31
32 #Layer 6
33 model3.add(Conv2D(36, (6, 6), padding='same', activation='relu'))
34 model3.add(BatchNormalization())
35 model3.add(Dropout(0.5))
36
37 #Layer 7
38 model3.add(Conv2D(24, (6, 6), padding='same', activation='relu'))
39 model3.add(BatchNormalization())
40 model3.add(MaxPooling2D(pool_size=(2, 2), padding='same'))
41 model3.add(Dropout(0.5))
42
43 model3.add(Flatten())
44 model3.add(Dense(128, activation='relu'))
45 model3.add(Dropout(0.5))
46 model3.add(Dense(num_classes, activation='softmax'))
47
48 print(model3.summary())
49
50 model3.compile(loss=keras.losses.categorical_crossentropy,
51 optimizer=keras.optimizers.Adadelta(),
52 metrics=['accuracy'])
53
54 history = model3.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
55
```



Model: "sequential_12"

Layer (type)	Output Shape	Param #
conv2d_63 (Conv2D)	(None, 28, 28, 70)	2590
batch_normalization_49 (Batch Normalization)	(None, 28, 28, 70)	280
conv2d_64 (Conv2D)	(None, 28, 28, 80)	201680
batch_normalization_50 (Batch Normalization)	(None, 28, 28, 80)	320
max_pooling2d_52 (MaxPooling2D)	(None, 14, 14, 80)	0
dropout_63 (Dropout)	(None, 14, 14, 80)	0
conv2d_65 (Conv2D)	(None, 14, 14, 60)	172860
batch_normalization_51 (Batch Normalization)	(None, 14, 14, 60)	240
dropout_64 (Dropout)	(None, 14, 14, 60)	0
conv2d_66 (Conv2D)	(None, 14, 14, 50)	108050
batch_normalization_52 (Batch Normalization)	(None, 14, 14, 50)	200
max_pooling2d_53 (MaxPooling2D)	(None, 7, 7, 50)	0
dropout_65 (Dropout)	(None, 7, 7, 50)	0
conv2d_67 (Conv2D)	(None, 7, 7, 48)	86448
batch_normalization_53 (Batch Normalization)	(None, 7, 7, 48)	192
dropout_66 (Dropout)	(None, 7, 7, 48)	0
conv2d_68 (Conv2D)	(None, 7, 7, 36)	62244
batch_normalization_54 (Batch Normalization)	(None, 7, 7, 36)	144
dropout_67 (Dropout)	(None, 7, 7, 36)	0
conv2d_69 (Conv2D)	(None, 7, 7, 24)	31128
batch_normalization_55 (Batch Normalization)	(None, 7, 7, 24)	96
max_pooling2d_54 (MaxPooling2D)	(None, 4, 4, 24)	0
dropout_68 (Dropout)	(None, 4, 4, 24)	0
flatten_11 (Flatten)	(None, 384)	0
dense_21 (Dense)	(None, 128)	49280
dropout_69 (Dropout)	(None, 128)	0
dense_22 (Dense)	(None, 10)	1290
Total params: 717,042		
Trainable params: 716,306		
Non-trainable params: 736		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 23s 384us/step - loss: 1.0074 - acc: 0.6573 - val_loss: 0.1724 - val_acc: 0.96

Epoch 2/10

60000/60000 [=====] - 18s 296us/step - loss: 0.1268 - acc: 0.9660 - val_loss: 0.0655 - val_acc: 0.98

Epoch 3/10

60000/60000 [=====] - 18s 299us/step - loss: 0.0793 - acc: 0.9792 - val_loss: 0.0294 - val_acc: 0.99

Epoch 4/10

60000/60000 [=====] - 18s 297us/step - loss: 0.0627 - acc: 0.9842 - val_loss: 0.0299 - val_acc: 0.99

Epoch 5/10

60000/60000 [=====] - 18s 297us/step - loss: 0.0537 - acc: 0.9861 - val_loss: 0.0431 - val_acc: 0.98

Epoch 6/10

60000/60000 [=====] - 18s 297us/step - loss: 0.0486 - acc: 0.9877 - val_loss: 0.0221 - val_acc: 0.99

Epoch 7/10

60000/60000 [=====] - 18s 296us/step - loss: 0.0413 - acc: 0.9892 - val_loss: 0.0283 - val_acc: 0.99

Epoch 8/10

60000/60000 [=====] - 18s 296us/step - loss: 0.0399 - acc: 0.9904 - val_loss: 0.0281 - val_acc: 0.99

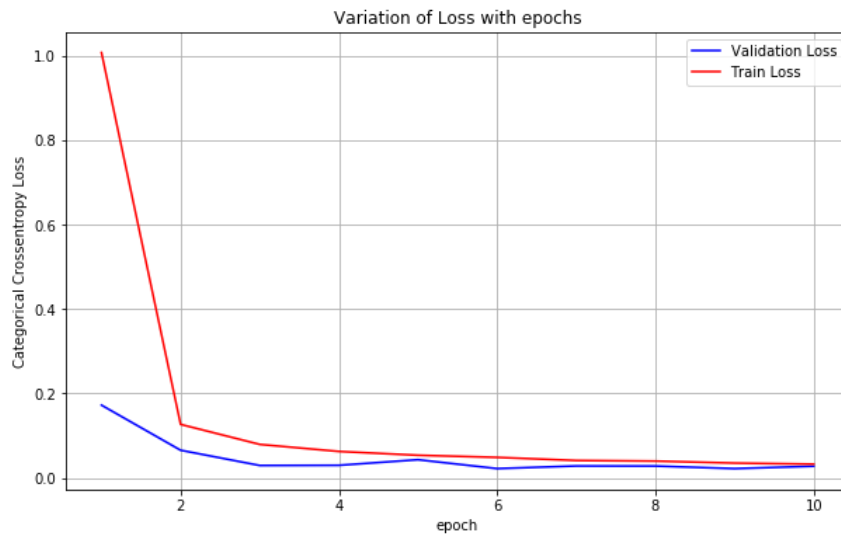
Epoch 9/10

60000/60000 [=====] - 18s 297us/step - loss: 0.0351 - acc: 0.9909 - val_loss: 0.0220 - val_acc: 0.99

Epoch 10/10

60000/60000 [=====] - 18s 298us/step - loss: 0.0326 - acc: 0.9914 - val_loss: 0.0277 - val_acc: 0.99

Test score: 0.02765477129972496
Test accuracy: 0.9945



Summary

```
1 # Please compare all your models using Prettytable library
2 from prettytable import PrettyTable
3 x = PrettyTable()
4 print('\nResults : CNN on MNIST data')
5 print(40*'=')
6
7 x.field_names = ["Hidden layers", "Kernel size", "Padding", "Max pooling", "epochs", "Train error", "Train accuracy", "Test error", "Test accuracy"]
8 x.add_row([3, '3X3', 'valid', '(2,2)', 10, 0.0623, '98.18 %', 0.0207, '99.35 %'])
9 x.add_row([5, '5X5', 'same', '(2,2)', 10, 0.0511, '98.60 %', 0.0194, '99.40 %'])
10 x.add_row([7, '6X6', 'same', '(2,2)', 10, 0.0326, '99.14 %', 0.0277, '99.45 %'])
11
12 print(x)
```

Results : CNN on MNIST data

Hidden layers	Kernel size	Padding	Max pooling	epochs	Train error	Train accuracy	Test error	Test accuracy
3	3X3	valid	(2,2)	10	0.0623	98.18 %	0.0207	99.35 %
5	5X5	same	(2,2)	10	0.0511	98.60 %	0.0194	99.40 %
7	6X6	same	(2,2)	10	0.0326	99.14 %	0.0277	99.45 %